

# Gradient Boosting Tree

Luu Minh Sao Khue

December 7, 2025

## 1. Introduction

Gradient Boosting Trees (GBT) are a powerful ensemble learning method that builds a strong predictive model by combining many weak decision trees. The main ideas are:

- Instead of training all trees independently (as in Random Forest), GBT builds trees *sequentially*.
- Each new tree tries to correct the errors made by the previous trees.
- The corrections are not done randomly: they follow the *gradient* of a chosen loss function.
- By repeatedly adding these small improvements, the final model becomes highly accurate and flexible.

Gradient boosting is general: it works for regression, binary classification, and multi-class classification simply by choosing an appropriate loss function.

## 2. Intuition

Suppose we want a model  $F(x)$  that minimizes some loss function  $L(y, F(x))$ . In standard boosting (e.g., AdaBoost), each new tree focuses on the “hard” samples. In contrast, *gradient boosting* uses a more principled view:

*Each new tree is trained to predict the **negative gradient** of the loss function with respect to the current predictions.*

**Why the gradient?** Because the gradient tells us the direction in which the loss decreases the fastest. If our model  $F_{m-1}(x)$  is not perfect, then the gradient of the loss at each training sample tells us how we should adjust the prediction at that point. A decision tree is then trained to approximate these gradient values.

Thus, the gradient appears as:

$$g_{im} = - \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \Big|_{F(x)=F_{m-1}(x)}.$$

This  $g_{im}$  becomes the pseudo-target that the  $m$ -th decision tree tries to fit.

### 3. Loss Functions and Derivative Derivations

In gradient boosting, each new tree is trained to approximate the *negative gradient* of the loss with respect to the current model prediction  $F(x)$ . Therefore, we must be able to **compute the derivative of the loss with respect to  $F(x)$** . This derivative tells us:

- how much the loss would change if we slightly changed the prediction  $F(x)$ ,
- in which direction we should move  $F(x)$  to reduce the loss.

The negative gradient at each training point becomes the *pseudo-target* that the next tree tries to fit.

Below we derive these gradients step by step.

#### 3.1 Regression Loss and Its Derivative

For standard regression, we often use the squared error loss:

$$L(y, F(x)) = \frac{1}{2}(y - F(x))^2.$$

Let us denote  $F(x)$  by  $F$  for simplicity. Then

$$L(y, F) = \frac{1}{2}(y - F)^2.$$

We want  $\frac{\partial L}{\partial F}$ .

**Step 1: Expand the expression conceptually.**

$$L = \frac{1}{2}(y - F)^2.$$

**Step 2: Apply the chain rule.** Treat  $(y - F)$  as an inner function:

$$u(F) = y - F, \quad L = \frac{1}{2}u^2.$$

Then:

$$\frac{dL}{dF} = \frac{dL}{du} \cdot \frac{du}{dF}.$$

**Step 3: Differentiate each part.**

$$\frac{dL}{du} = \frac{1}{2} \cdot 2u = u = y - F,$$

$$\frac{du}{dF} = \frac{d}{dF}(y - F) = -1.$$

**Step 4: Multiply them.**

$$\frac{\partial L}{\partial F} = (y - F) \cdot (-1) = F - y.$$

So the gradient with respect to the prediction is

$$\frac{\partial L}{\partial F(x)} = F(x) - y.$$

The **negative gradient** (the value we actually fit) is

$$-\frac{\partial L}{\partial F(x)} = y - F(x),$$

which is exactly the usual regression residual.

### 3.2 Binary Classification (Logistic Loss) and Its Derivative

For binary classification, we often use labels  $y \in \{-1, +1\}$  and the logistic loss

$$L(y, F(x)) = \log(1 + e^{-yF(x)}).$$

Again, write  $F(x)$  as  $F$  to simplify notation:

$$L(y, F) = \log(1 + e^{-yF}).$$

**Step 1: Apply the chain rule.** Let

$$u(F) = 1 + e^{-yF}, \quad L = \log u.$$

Then

$$\frac{dL}{dF} = \frac{dL}{du} \cdot \frac{du}{dF}.$$

**Step 2: Differentiate each part.** First:

$$\frac{dL}{du} = \frac{1}{u} = \frac{1}{1 + e^{-yF}}.$$

Next, for  $u(F) = 1 + e^{-yF}$ :

$$\frac{du}{dF} = \frac{d}{dF}(1 + e^{-yF}) = 0 + e^{-yF} \cdot \frac{d(-yF)}{dF} = e^{-yF} \cdot (-y) = -ye^{-yF}.$$

**Step 3: Multiply them.**

$$\frac{\partial L}{\partial F} = \frac{1}{1 + e^{-yF}} \cdot (-ye^{-yF}) = -y \frac{e^{-yF}}{1 + e^{-yF}}.$$

Notice that

$$\frac{e^{-yF}}{1 + e^{-yF}} = \frac{1}{e^{yF} + 1} = \frac{1}{1 + e^{yF}},$$

so we can rewrite:

$$\frac{\partial L}{\partial F} = -\frac{y}{1 + e^{yF}}.$$

The **negative gradient** used as the pseudo-target is then

$$-\frac{\partial L}{\partial F(x)} = \frac{y}{1 + e^{yF(x)}}.$$

**Logit and Probability: Why This Form Appears** In gradient boosting for classification, the model output  $F(x)$  is often interpreted as a *logit*, i.e. a log-odds:

$$F(x) = \log \frac{p(x)}{1 - p(x)},$$

where  $p(x) = \mathbb{P}(y = +1 | x)$ .

We can convert from logit  $F$  to probability  $p$  as follows:

$$F = \log \frac{p}{1 - p} \implies e^F = \frac{p}{1 - p}.$$

Solve for  $p$ :

$$\begin{aligned} e^F(1 - p) &= p \implies e^F - e^F p = p \\ e^F &= p + e^F p = p(1 + e^F) \\ p &= \frac{e^F}{1 + e^F} = \frac{1}{1 + e^{-F}}. \end{aligned}$$

Thus the probability  $\mathbb{P}(y = +1 | x)$  is

$$p(x) = \sigma(F(x)) = \frac{1}{1 + e^{-F(x)}}.$$

Similarly,

$$\mathbb{P}(y = -1 | x) = 1 - p(x) = \frac{1}{1 + e^{F(x)}}.$$

This is exactly the term that appears in the gradient:

$$\frac{\partial L}{\partial F} = -\frac{y}{1 + e^{yF}} = \begin{cases} -\frac{1}{1 + e^F} & \text{if } y = +1, \\ +\frac{1}{1 + e^{-F}} & \text{if } y = -1, \end{cases}$$

which can be expressed in terms of probabilities of the classes.

## 4. Binary Classification (Logistic Loss): Logit-Space and Probability-Space Gradients

For binary classification with the logistic loss, there are **two equivalent ways** to compute the gradient. They differ only in the representation of the labels and the model output, but both lead the boosting algorithm to update predictions in the same direction.

**Method 1: Logit-Space Formulation (Original Boosting View).** We use labels  $y \in \{-1, +1\}$  and the model output  $F(x)$  is a *logit* (log-odds). The logistic loss is:

$$L(y, F) = \log(1 + e^{-yF}).$$

As derived earlier, the gradient is:

$$\frac{\partial L}{\partial F} = -\frac{y}{1 + e^{yF}},$$

and the **negative gradient** (the residual used as a pseudo-target) is:

$$r_i^{(\text{logit})} = \frac{y_i}{1 + e^{y_i F(x_i)}}.$$

This is the formulation used in the original gradient boosting framework.

**Method 2: Probability-Space Formulation.** The model can be rewritten in terms of a **probability**:

$$p_i = \sigma(F(x_i)) = \frac{1}{1 + e^{-F(x_i)}},$$

and convert labels to  $y'_i \in \{0, 1\}$ .

In probability space, the gradient of the logistic (cross-entropy) loss w.r.t. the probability is:

$$r_i^{(\text{prob})} = y'_i - p_i.$$

This looks intuitive: *residual = actual class – predicted probability*. However, boosting still needs to update the *logit*  $F(x)$ , not the probability. Therefore, using this method must convert:

$$F = \log \frac{p}{1-p}, \quad p = \sigma(F).$$

Thus, after fitting a tree to  $y' - p$ , the update is applied to the logit  $F(x)$ , and then optionally converted back to a probability.

**Why the Two Methods Are Equivalent.** The key fact is that both gradients point in the *same improvement direction*. If  $y' = 1$  but  $p$  is too small, both gradients tell the model to increase  $F(x)$ ; if  $y' = 0$  but  $p$  is too large, both gradients tell the model to decrease  $F(x)$ .

Mathematically, the two expressions are linked by the identity:

$$\frac{y}{1 + e^{yF}} = \begin{cases} 1 - p & \text{if } y' = 1, \\ -p & \text{if } y' = 0, \end{cases}$$

where  $y' = (y + 1)/2$  converts  $\{-1, +1\}$  labels to  $\{0, 1\}$ .

Thus:

$$r_i^{(\text{logit})} \propto y'_i - p_i = r_i^{(\text{prob})}.$$

The two forms differ by a positive scaling factor, which does not affect the direction of the fitted tree. Therefore, both methods are theoretically consistent and lead to the same gradient-boosting updates.

## 5. Mathematical Algorithm

Let  $(x_i, y_i)$ ,  $i = 1, \dots, N$ , be the training data, and let  $M$  be the total number of trees. Let  $\nu$  be the learning rate.

---

**Algorithm 1** Gradient Boosting Tree (General Form)

---

1: **Initialize** the model with a constant

$$F_0(x) = \arg \min_c \sum_{i=1}^N L(y_i, c).$$

2: **for**  $m = 1$  **to**  $M$  **do**

3:   Compute negative gradients (pseudo-residuals):

$$g_{im} = - \left. \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right|_{F(x)=F_{m-1}(x)}, \quad i = 1, \dots, N.$$

4:   Fit a regression tree  $h_m(x)$  to the targets  $\{g_{im}\}_{i=1}^N$ .

5:   For each leaf  $j$  of the tree, compute the optimal leaf weight

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma),$$

where  $R_{jm}$  is the set of samples in leaf  $j$  of tree  $h_m$ .

6:   Update the model:

$$F_m(x) = F_{m-1}(x) + \nu \sum_j \gamma_{jm} \mathbb{I}(x \in R_{jm}).$$

7: **end for**

8: **Return** the final model  $F_M(x)$ .

---

## 6. Summary

Gradient boosting builds a strong model by:

- Starting from a simple constant prediction.
- Repeatedly fitting decision trees to the negative gradient of the loss.
- Updating the model slowly using a learning rate.

This framework unifies regression, classification, and many other tasks, and forms the foundation for modern algorithms such as XGBoost, LightGBM, and CatBoost.